

Phonorm: Phonetic Text Normalization using Recurrent Neural Networks

Jasper Ginn, Arvid Halma, Aaron hislastname & Boaz Manger

December 18, 2018

1 Introduction

Text normalization is the application of linguistic models to text with the aim of standardizing words, which are also referred to as **tokens**. For example, tokens may refer to the same underlying concepts but can be spelled in a different way. Hence, the aim of text normalization is to reduce tokens to a common form. Examples of normalization are **stemming** and **lemmatization**. The former tries to map words to tokens by removing the suffix of a word (e.g. ‘economy’, ‘economist’ and ‘economic’ all map to ‘econom’), and the latter approach attempts to bring back each token to its dictionary representation.

Phonetic normalization attempts to reduce words that sound similar to the same token (Bisani and Ney 2008). Due to the proliferation of social media, there has been a surge of interest in this type of normalization. Text found on social media such as Twitter, which is often referred to as *microtext*, has different challenges than regular, well-written text. It is characterized by a casual writing style, colloquial tone, relaxed spelling, reliance on emoticons, use of out of vocabulary (OOV) words such as acronyms and

emphasis on emotional undercurrents for dramatic effect. (See e.g. Doval, Vilares, and Vilares 2018, Xue, Yin, and Davison (2011), López-Ludeña et al. (2012), Li and Liu (2012), Satapathy et al. (2017), Khoury (2015b), Khoury (2015a)).

Our interest in this topic is related to the development of our chatbot **ChitChat** (Halma 2019), which we have implemented in for several international organizations in countries in which English is not the first language. One of our findings is that many of the conversations are characterized by a use of English that strongly resembles its phonetic form. Furthermore, while we find that the language used to converse with our chatbot bears some similarities to microtext, the language used is often more formal. This is an important distinction because strict phonetic normalization as we implement it is not suited for the informal style of microtext (see Doval, Vilares, and Vilares 2018).

In this paper, we present an approach to phonetic normalization that employs a standard encoder-decoder framework common to machine translation tasks. In this setting, the inputs are characters of English words, and the outputs are the pronunciations of these words in various formats. Hence, the goal of this paper is to show that a mapping can be

learned from written to phonetic English by using character-level Recurrent Neural Networks (RNNs). Because RNNs excel at learning sequential patterns, they are a natural choice when considering the problem of phonetic normalization. RNNs have shown promise in diverse applications of Natural-Language Processing (NLP). For example, X and Y showed that RNNs do well in translation tasks. One benefit of such an approach is that the mapping is learned without specific expertise about such a task. Examples of previous work with respect to text normalization and deep learning are e.g. Min and Mott (2015), Chrupala (2019) and Jaitly and Sproat (2017)].

This paper is structured as follows. The first section gives a short description of RNNs. We then outline the data sets, models and evaluation metrics used in our experiment. Section three interprets the results. We close by offering several ways to improve the model. We use two publicly available datasets and have made all the code and data available on GitHub.¹

2 Recurrent neural networks

Recurrent neural networks (RNN) are a type of neural network that model sequences such as series of words or characters. The RNN takes both an input vector, such as a sequence of input characters, as well as outputs that occur earlier in the sequence. Therefore, at timestep t , the results of timestep $t - 1$ are fed into the network to learn patterns, or sequences,

¹You can find the repository at <https://github.com/JasperHG90/phonorm>

that are inherent to the data.

In principle, the length of the sequence is arbitrary, and the input length T_x does not have to match the output length T_y . This is convenient for language modeling because, for example, an English word rarely has the same length as its counterpart in another language.

Like other models that consist of many layers, RNNs suffer from the ‘vanishing gradient’ problem. This means that they cannot represent long-range dependencies very well, and that, for example, a verb that occurs early in the sentence may not be ‘memorized’ until the end of the sentence, which renders the RNN unable to establish the connection between the verb and a noun or adjective that refers to it. This problem is addressed by making use of RNNs that are able to capture long-range dependencies such as Long Short-Term Memory (LSTM) layers or Gated Recurrent Units (GRU). (Hochreiter and Schmidhuber 1997, Cho et al. (2014)). In either one of these models, the model can be *unidirectional* or *bidirectional*. In the latter case, the model uses information from both earlier as well as later in the sequence to learn patterns (Schuster and Paliwal 1997).

3 Data and models

This section describes the data used to train the models, and how the models are configured.

3.1 Data

We consider two separate datasets that are subdivided into three categories:

3.1.1 cmudict

a dataset containing 133.031 words mapping English words to their phonetic representation using ARPAbet phoneme set (“CMU Pronouncing Dictionary” 2018). We remove words containing digits, special punctuation and special characters. The final dataset contains 123.612 words. The data are subdivided into two data sets:

- *Multiple*: treats phonemes as individual characters. That is, a word like ‘abate’ is translated as ‘AA B EY T’ and would therefore contain the vocabulary entries ‘AA’, ‘B’, ‘EY’ and ‘T’.
- *Single*: the special arrangement between letters in the reference are ignored. Hence, the same word (‘abate’) now contains the vocabulary entries ‘A’, ‘B’, ‘T’, ‘E’, ‘Y’.

This distinction between the datasets allow us to observe whether it is useful to retain the phonemes or not.

3.1.2 Wiktionary

This is a dataset that is retrieved from the wiktionary dump file (Wikipedia 2018). We used the ‘wikt2pron’ python module (Xiong 2018) to extract all English words from the dump from 01–11–2018.

The data contains 46.675 words with pronunciations in XSAMPA format after removal of entries containing digits and special characters.

One downside of the wikt2pron module (or from the wiktionary data set) is that it does not distinguish between dialects. That is, the module returns a python dictionary for each word as well as up to 20 pronunciation entries without specifying whether these are British-English, American-English or, for example, Australian-English. Clearly, this is an issue because British homophones do not map to American homophones one-to-one.

3.2 Evaluation metrics

We evaluate the model on three separate metrics.

3.2.1 BLEU score

The BLEU score (Papineni et al. 2002) is a measure of overlap between the predicted word and a reference that outputs a score between 0 – 1, where a score of 0 means that there is no overlap and 1 means that there is perfect overlap.

Typically, the BLEU score is a combined metric measuring the overlap of one-grams, two-grams, three-grams and four-grams, each of which are weighted equally. Moreover, it takes into account the length of words (since shorter words are easier to predict). In this paper, we use the `sentence_bleu` function from the `nltk` python module (Bird, Klein, and Loper 2009). We only present the cumulative 4-gram score (BLEU-4).

3.2.2 Accuracy in detecting homophones

A natural way to test the accuracy of the model is to evaluate its performance on homophones. These are words that are spelled differently but that are pronounced the same. We define three metrics to measure the performance of our models in detecting homophones, all of which are proportions p . These three metrics are as follows:

1. **all equal**: the proportion of observations for which all homophones map to the same pronunciation.
2. **equal to reference**: the proportion of observations for which all homophones are equal to the reference.
3. **largest subgroup**: the average of the largest sub-group of unique pronunciations for each example.

Let θ_i be the value of the i^{th} development set example. If all the predictions are equal, the value of $\theta_i = 1$, else, it equals 0, then

$$p_{\text{all equal}} = \frac{1}{n} \sum_{i=1}^n \theta_i \quad (1)$$

If we let θ_i be the value of the j^{th} word out of k words in the i^{th} development set example, which equals 1 if the j^{th} word equals the reference and 0 otherwise, then:

$$p_{\text{equal to reference}} = \frac{1}{n} \sum_{i=1}^n \left[\frac{1}{k} \sum_{j=1}^k \theta_i \right] \quad (2)$$

Finally, if we let u equal the largest subgroup of k possible subgroups, where k is the number of words in the i^{th} development set example. Furthermore, if we let θ_i be the value of the j^{th} word out of k words in the i^{th} development set example, which equals 1 if the j^{th} word equals the value of the largest subgroup u and 0 otherwise, then

$$p_{\text{largest subgroup}} = \frac{1}{n} \sum_{i=1}^n \left[\frac{1}{k} \sum_{j=1}^k \theta_i \right] \quad (3)$$

3.2.3 Accuracy in recognizing similar-sounding misspelled words

As a third evaluation of the model, we consider the problem of mapping misspelled - but phonetically similar - words to their phonetic representation. We take a random subset of 100 commonly misspelled words provided by Wikipedia (wikipedia 2018) and manually classify these words as either 1 if the misspelled word should be phonetically similar to the correctly spelled word and 0 otherwise. This dataset gives us an idea of the extent to which the models perform on unseen and unusual data. The metrics used to evaluate this data set are accuracy, sensitivity, specificity and the F1 score.

3.3 Cross-validation

The data are split into several train and test sets. We first split the data based on homophones. This process works as follows: we first extract all homophones from the dataset, after which we select a subset as the test set. All words that map to the same pronun-

ciation are then added to the test set. This ensures that the pronunciation is unseen during training.

To evaluate the general skill (BLEU score) of the model, we then split the train set into a generic train and test split.

3.4 Models

This section describes the architecture of our models, all of which are trained using the python library *Keras* (Chollet 2015).

3.4.1 Hyperparameters and training

All models are trained using stochastic gradient descent and Adam (Kingma and Ba 2014) with mini-batch size 128. The optimization metric is the *log-likelihood*. We apply dropout and recurrent dropout (Gal and Ghahramani 2015) with 0.1 to both the encoder and decoder. We also use ‘teacher forcing’, meaning that instead of ingesting the predicted character generated at timestep $t - 1$ to predict the next character in the sequence, we use the ground truth character (Lamb et al. 2016). All models are trained for 10 epochs on a single NVIDIA GTX 1060 with 6GB RAM.

Table 1: Hyperparameters for final models

	batch size	epochs	validation split	hidden dimension	learning rate
cmudict (multiple)	128	10	0.05	512	0.005
cmudict (single)	128	10	0.05	512	0.005
wiktionary	128	10	0.05	512	0.002

3.4.2 Encoder and decoder

The English words and their phonetic representation in our dataset are first converted to a one-hot-

encoding (OHE). We then use a bidirectional LSTM that takes the OHE English words as input and that passes its internal state to the decoder. The decoder takes as its input the state vectors from the encoder as well as the OHE phonetic words. The decoder then generates the output sequence.




Figure 1: Encoder/decoder architecture

4 Results

This section describes the results of the models. Figure 2 below shows the average BLEU score for words of different lengths.

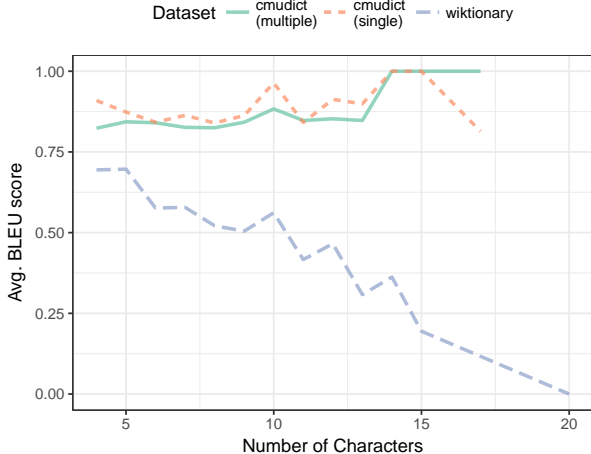


Figure 2: Average BLEU score on development set

The models trained on the *cmudict* data vastly outperform the *wiktionary* data. The *cmudict*

(single) model achieves an average BLEU score of 86.78%, which is slightly higher than the `cmudict` (multiple) model at 83.37%.

Table 1 below shows the performance of the models on the homophone test set. The `cmudict` (single) model outperforms the other models on all metrics. However, it is interesting that the models are much closer in terms of performance on the most important metric ($p_{\text{largest subgroup}}$) than on the other metrics. It should be noted that these metrics do not measure any possible erroneous mappings (‘false positives’). Hence, it may be the case that the `wiktionary` model achieves this high score at the cost of making many more mistakes than the `cmudict` data.

Table 2: Percentage correct on each metric

	All equal	Largest subgroup	Equal to reference
<code>cmudict</code> (single)	77.33%	91.15%	86.94%
<code>cmudict</code> (multiple)	63.56%	86.58%	76.87%
<code>wiktionary</code>	42.11%	80.2%	51.02%

The accuracy across word length is similar across datasets, indicating that they perform well irrespective of the size of a word.

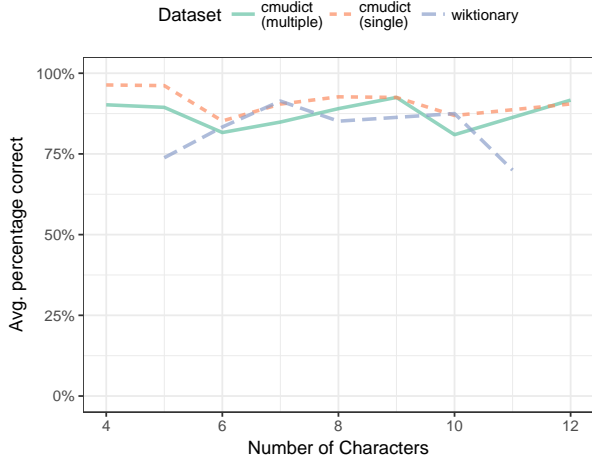


Figure 3: Average percentage correct (largest subgroup) across word length

Finally, we consider the performance of the model on misspelled words. Given that roughly one-third of our manually labelled data is considered to be of the positive class², we are most interested in the F1 and sensitivity scores obtained by the model.

Table 3: Evaluation on misspelled data set

	<code>cmudict</code> (single)	<code>cmudict</code> (multiple)	<code>wiktionary</code>
Accuracy	0.94	0.83	0.76
Sensitivity	0.87	0.77	0.69
Specificity	0.92	0.89	0.87
F1	0.90	0.83	0.77

As can be observed from table 2, the `cmudict` (single) model again outperforms the others. It achieves significantly better scores than even the `cmudict` (multiple) model.

5 Discussion and conclusion

The RNN approach to phonetic normalization is, as is shown above, successful in its main objective.

²indicating that the pronunciation of the misspelled word is the same as the word that was intended

That is, the models succeed in creating sensible output, are good at mapping homophones to the same phonetic representation, and perform adequately on misspelled words. On all metrics, the `cmudict (single)` model outperforms the other models. This is a rather curious outcome, as one would expect the model that uses graphemes (`cmudict (multiple)`) to incorporate more information and hence perform better.

There are several improvements to be made to the model. Firstly, there must be some noise in the data because it is not-context aware. This means that the model cannot learn the difference between, for example, a word like ‘tear’ (to tear a piece of paper) or ‘tear’ (he teared up). This type of word is called a ‘heteronym’. Adding this kind of context awareness must be gauged from the sentence somehow, given that the models described in this paper deal exclusively with words.

Furthermore, while the model performed adequately on the wikipedia data, it faces several bottlenecks. Firstly, the size of the vocabulary is not ideal. This issue could be exacerbated for languages other than English, as we expect the size of vocabularies to be smaller for other languages that are not as widely used on the internet. Secondly, more work should be done to filter different dialects or pronunciations from this data. As mentioned earlier, this project used available software to extract words from the wiktionary. This software does not distinguish between different pronunciations and hence we were not able to filter for, for example, American or

British pronunciations.

Finally, we have implemented a rather traditional infrastructure. RNNs are resource-heavy from an architectural point of view. Other models (for example pure Attention models or models that employ convolutional layers) have been shown to perform as good or better while substantially reducing the resources needed to train the model.

References

- Bird, Steven, Ewan Klein, and Edward Loper. 2009. *Natural Language Processing with Python*.
- Bisani, Maximilian, and Hermann Ney. 2008. “Joint-Sequence Models for Grapheme-to-Phoneme Conversion.” *Speech Communication* 50 (5): 434–51. doi:10.1016/j.specom.2008.01.002.
- Cho, Kyunghyun, Bart van Merriënboer, Caglar Gulcehre, Dzmitry Bahdanau, Fethi Bougares, Holger Schwenk, and Yoshua Bengio. 2014. “Learning Phrase Representations Using RNN Encoder-Decoder for Statistical Machine Translation.” *arXiv:1406.1078 [Cs, Stat]*, June. <http://arxiv.org/abs/1406.1078>.
- Chollet, François. 2015. “Keras.” <https://keras.io>.
- Chrupała, Grzegorz. 2019. “Normalizing Tweets with Edit Scripts and Recurrent Neural Embeddings.” Accessed January 3. <http://citeseerx.ist.psu.edu/viewdoc/citations;jsessionid=3C665FAF5D29C92C6D8CB06DFF03EC79?doi=>

- 10.1.1.652.9864.
- “CMU Pronouncing Dictionary.” 2018. *Wikipedia*. https://en.wikipedia.org/w/index.php?title=CMU_Pronouncing_Dictionary&oldid=871723674.
- Doval, Yeraí, Manuel Vilares, and Jesús Vilares. 2018. “On the Performance of Phonetic Algorithms in Microtext Normalization.” *Expert Systems with Applications* 113 (December): 213–22. doi:10.1016/j.eswa.2018.07.016.
- Gal, Yarin, and Zoubin Ghahramani. 2015. “A Theoretically Grounded Application of Dropout in Recurrent Neural Networks.” *arXiv:1512.05287 [Stat]*, December. <http://arxiv.org/abs/1512.05287>.
- Halma, Arvid. 2019. “Arvid / Chitchat.” Accessed January 3. <https://bitbucket.org/arvid/chitchat>.
- Hochreiter, Sepp, and Jürgen Schmidhuber. 1997. “Long Short-Term Memory.” *Neural Comput.* 9 (8): 1735–80. doi:10.1162/neco.1997.9.8.1735.
- Jaitly, Navdeep, and Richard Sproat. 2017. “An RNN Model of Text Normalization.” In. <https://arxiv.org/pdf/1611.00068.pdf>.
- Khoury, R. 2015a. “Phonetic Normalization of Microtext.” In *2015 IEEE/ACM International Conference on Advances in Social Networks Analysis and Mining (ASONAM)*, 1600–1601. doi:10.1145/2808797.2809352.
- . 2015b. “Microtext Normalization Using Probably-Phonetically-Similar Word Discovery.” In *2015 IEEE 11th International Conference on Wireless and Mobile Computing, Net-*
- working and Communications (WiMob)*, 384–91. doi:10.1109/WiMOB.2015.7347988.
- Kingma, Diederik P., and Jimmy Ba. 2014. “Adam: A Method for Stochastic Optimization.” *arXiv:1412.6980 [Cs]*, December. <http://arxiv.org/abs/1412.6980>.
- Lamb, Alex, Anirudh Goyal, Ying Zhang, Saizheng Zhang, Aaron Courville, and Yoshua Bengio. 2016. “Professor Forcing: A New Algorithm for Training Recurrent Networks.” *arXiv:1610.09038 [Cs, Stat]*, October. <http://arxiv.org/abs/1610.09038>.
- Li, Chen, and Yang Liu. 2012. “Normalization of Text Messages Using Character- and Phone-Based Machine Translation Approaches.” In. Vol. 3.
- López-Ludeña, Verónica, Rubén San-Segundo, J. M. Montero, Roberto Barra-Chicote, and Javier Mateos Lorenzo. 2012. “Architecture for Text Normalization Using Statistical Machine Translation Techniques.” In.
- Min, Wookhee, and Bradford W. Mott. 2015. “NCSU_SAS_WOOKHEE: A Deep Contextual Long-Short Term Memory Model for Text Normalization.” In *NUT@IJCNLP*. doi:10.18653/v1/W15-4317.
- Papineni, Kishore, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. “BLEU: A Method for Automatic Evaluation of Machine Translation.” In *Proceedings of the 40th Annual Meeting on Association for Computational Linguistics*, 311–18. ACL ’02. Stroudsburg, PA, USA: Association for Computa-

tional Linguistics. doi:10.3115/1073083.1073135.

Satapathy, R., C. Guerreiro, I. Chaturvedi, and E. Cambria. 2017. “Phonetic-Based Micro-text Normalization for Twitter Sentiment Analysis.” In *2017 IEEE International Conference on Data Mining Workshops (ICDMW)*, 407–13. doi:10.1109/ICDMW.2017.59.

Schuster, M., and K. K. Paliwal. 1997. “Bidirectional Recurrent Neural Networks.” *IEEE Transactions on Signal Processing* 45 (11): 2673–81. doi:10.1109/78.650093.

Wikipedia. 2018. “Wikimedia Downloads.” Dataset. *Index of /Enwiktionary/*. Accessed December 28. <https://dumps.wikimedia.org/enwiktionary/20181101/enwiktionary-20181101-pages-meta-current.xml.bz2>.

wikipedia. 2018. “Wikipedia:Lists of Common Misspellings/for Machines.” *Wikipedia*. https://en.wikipedia.org/w/index.php?title=Wikipedia:Lists_of_common_misspellings/For_machines&oldid=838176804.

Xiong, Yifan. 2018. “Toolkit Converting Pronunciation in Enwiktionary Xml Dump to Cmuclt Format: Abuccts/Wikt2pron.” <https://github.com/abuccts/wikt2pron>.

Xue, Zhenzhen, Dawei Yin, and Brian Davison. 2011. “Normalizing Microtext.” In.